

Aishvarya Salvi

Buffalo, NY • +1 7164657631 • aishwarya.salvi28@gmail.com • [LinkedIn](#) • [GitHub](#) • [Portfolio](#)

Profile

AI/Data Engineer with 4+ years of experience building agentic LLM systems, RAG pipelines, and production ETL/ELT on Azure and AWS. Skilled in Python, SQL, LangChain/LangGraph, Snowflake, and FastAPI, with proven impact shipping explainable, auditable AI to non-technical stakeholders.

Education

University at Buffalo-NY , Master of Science in Data Science	2025
University at Pune-India , Bachelors in Computer Science	2021

Professional Experience

AI Engineer, [Research Foundation for SUNY](#) 01/25 – Present

- Architected a production agentic RAG system for supply chain disruption detection using LangChain ReAct with six domain-specific tools, generating real-time supplier risk scores and alternative-supplier recommendations for procurement teams.
- Engineered a semantic embedding pipeline with OpenAI text-embedding-3-small and a tiered Pinecone + pgvector store, persisting user-flagged alerts as a feedback loop to continuously improve agent retrieval and scoring.
- Built a multi-signal risk scoring engine combining VADER sentiment, LLM-as-judge severity, and recency decay into a calibrated 0–10 score, with content-hash caching eliminating redundant LLM calls.
- Shipped an async FastAPI backend with a Streamlit dashboard with agent guardrails, deployed via Docker Compose.

AI/ML Engineer, [New Era Cap](#) 01/25 – 02/26

- Architected an end-to-end demand forecasting platform on Azure Data Factory + Snowflake, replacing manual Excel planning and cutting stockout/overstock risk across multi-year horizons.
- Engineered quantile CatBoost models with lag, rolling, and seasonal features, achieving 87% accuracy and benchmarked against XGBoost, LightGBM, and SARIMA.
- Built a Snowflake ETL layer of CTEs and intermediate views over hundreds of raw tables with row-count and checksum data quality checks, delivering 10x faster analytical query performance.
- Integrated OpenAI via FastAPI into a Power BI + Streamlit stack, converting SHAP feature importance into plain-language explanations and shortening planning cycles by 40%.

Senior Data Analyst, [Mediaocean](#) 04/23 – 07/24

- Led data quality analysis by consolidating Python, Karate, Jenkins, and Docker defect data into structured datasets, lifting coverage from 60% to 95% and cutting escaped anomalies by 40%.
- Built Power BI defect dashboards for CTR, Actions, and Impressions and embedded SQL validation scripts into CI/CD pipelines with Agile teams.

Data Analyst, [Mediaocean](#) 08/21 – 03/23

- Validated digital ad campaign workflows using Python, SQL, and A/B tests across 450+ defects, cutting bug resolution time 55% and sustaining 97% build success.

Skills

- **Programming & Analytics:** Python, SQL, R, Pandas, NumPy, Scikit-learn, XGBoost, CatBoost, LightGBM
- **GenAI & Agentic Systems:** LangChain, LangGraph, RAG, ReAct Agents, OpenAI / Claude APIs, Pinecone, pgvector, sentence-transformers, LLM-as-judge, Pydantic guardrails
- **Data Engineering & Warehouses:** ETL/ELT, Snowflake, PostgreSQL, MySQL, Apache Airflow, FastAPI,
- **Cloud & Orchestration:** Azure (Data Factory, Blob Storage, Key Vault), AWS (S3, EC2, RDS, Redshift), Docker, Docker Compose, Jenkins, CI/CD, Git/GitHub
- **Analytics & BI Tools:** Power BI (DAX, Power Query, semantic models), Tableau, Streamlit, Matplotlib/Seaborn, SSE Streaming, Operational Reporting

Projects

[Multi-Agent Document Verification Platform](#), (LangGraph, RAG, FastAPI)

- Designed Built a five-agent LangGraph + RAG pipeline (Extractor, Validator, Cross-Reference, Compliance, Decision) that automates ID, contract, and certificate verification with auditor-grade reasoning traces, replacing manual review and achieving high precision on 500+ labeled documents via LLM-as-judge evaluation.

[LLM-Powered Utility Billing Audit Platform](#), (Airflow, LLM, PostgreSQL)

- Built an agentic LLM audit platform that parses 600+ pages of tariff documents into versioned SCD Type 2 rules and auto-detects billing anomalies with auditor-ready explanations, reducing manual audit time by 60% and scaling to 1,000+ bills per month.